

Tanzu GemFire

Building Faster Cloud Native
Applications at Scale with VMware
Tanzu GemFire

Getting Started

Download Source Code

<https://github.com/ggreen/spring-geode-showcase.git>



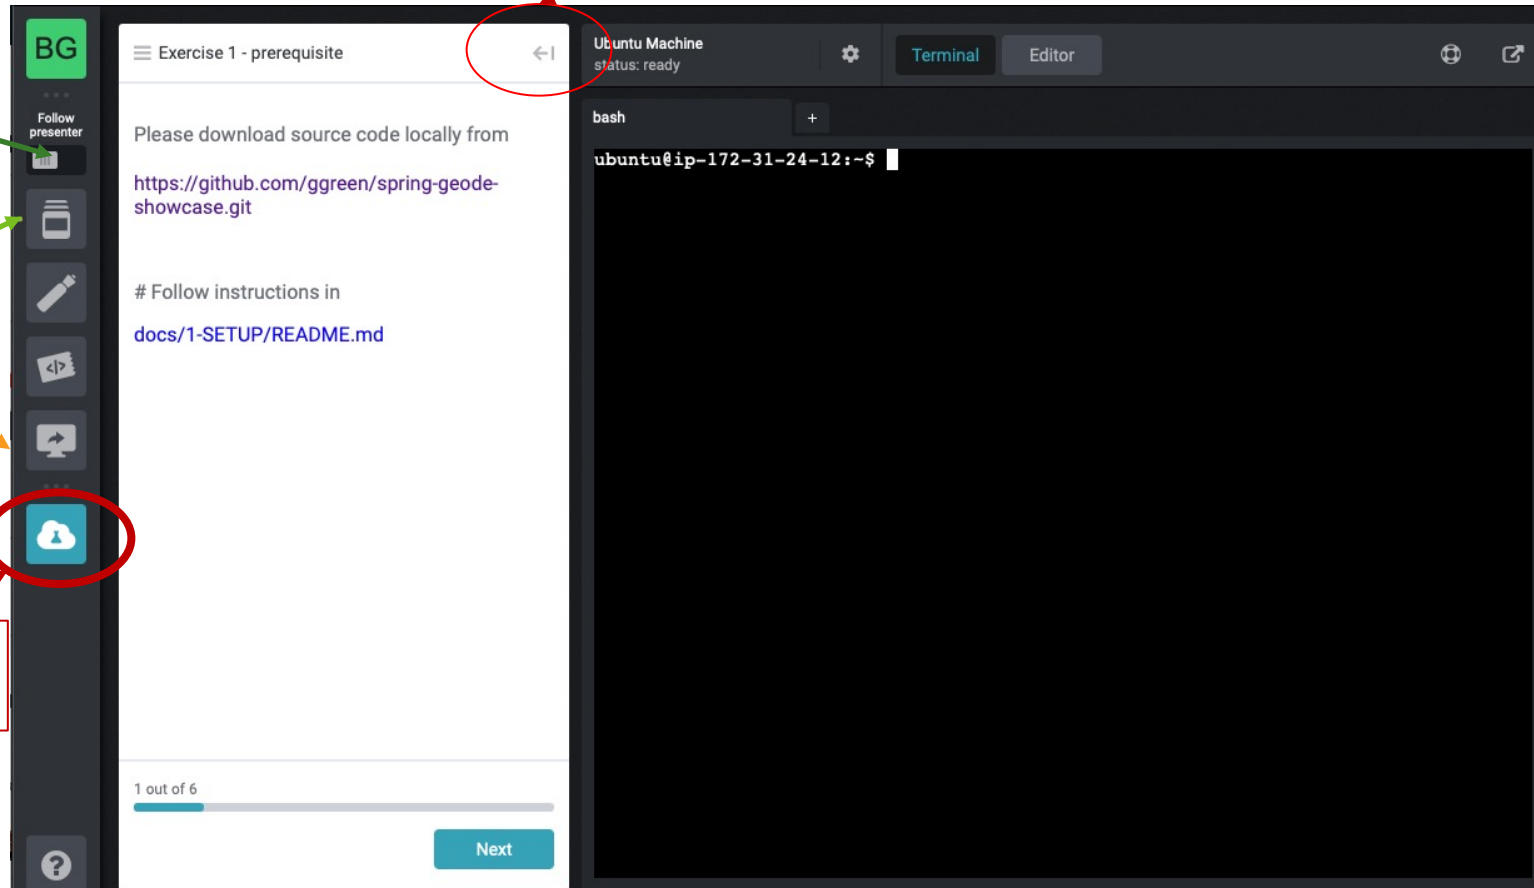
Click arrow If you do not see the exercises

Following presenter

See slides

See presenter screen

Click on My lab



Exercise 1 - prerequisite

Please download source code locally from

<https://github.com/ggreen/spring-geode-showcase.git>

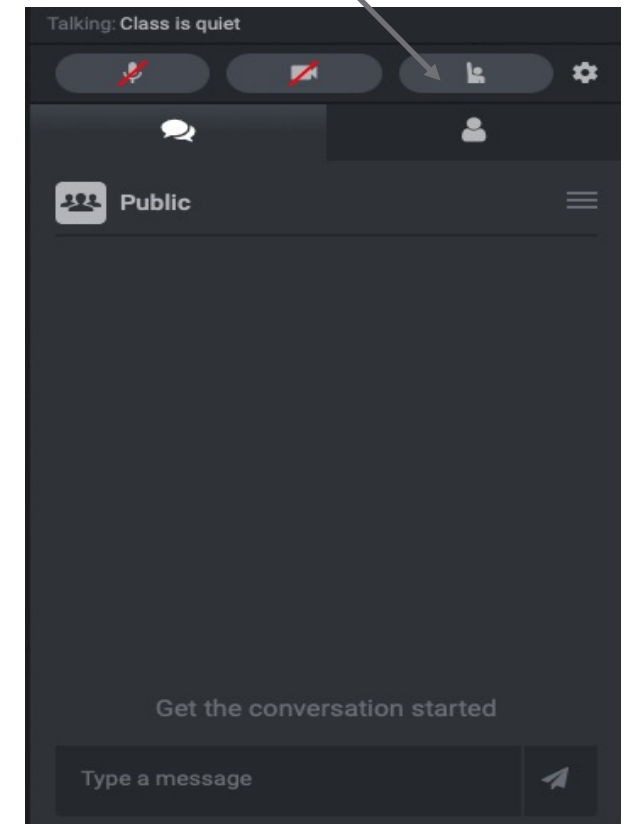
Follow instructions in

[docs/1-SETUP/README.md](#)

1 out of 6

Next

If you need assistance



Talking: Class is quiet

Public

Get the conversation started

Type a message

VMware Tanzu – Data Services



VMware Tanzu

Infrastructure for running modern apps and backing services with consistent, conformant Kubernetes everywhere.



Data Management
Management for
Tanzu Data
Services instances



GemFire

Fast In-Memory
data store for
Caching,
Transactional and
NoSQL support
powered by
Apache Geode

I need a
fast data
store



SQL

Relational MySQL
or Postgres
database for
Transactional or
Analytic data
processing

I need to
replatform a
relational
database



Greenplum

Massively Parallel
Processing (MPP)
Postgres for Big
Data store for
analytics, Machine
Learning and
Artificial Intelligence

I need to drive
analytic value
of out tons of
existing data



Rabbit MQ

High throughput broker for
reliable messaging delivery

I need reliable
messaging delivery



Spring Cloud Data Flow

Data integration
orchestration service for
dynamically building data
pipelines

I need flexible and
manageable data
integrations

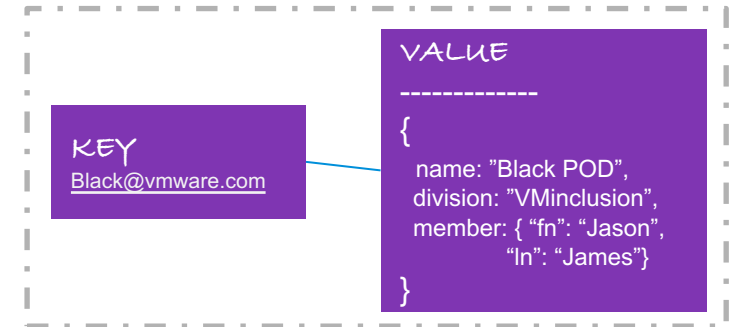
Features

- ✓ Cloud deployed backing-services
- ✓ On-Premise and Multi-Cloud
- ✓ Based on open source
- ✓ Scaling
- ✓ HA - Fault Tolerant
- ✓ Secured access
- ✓ World Class Support

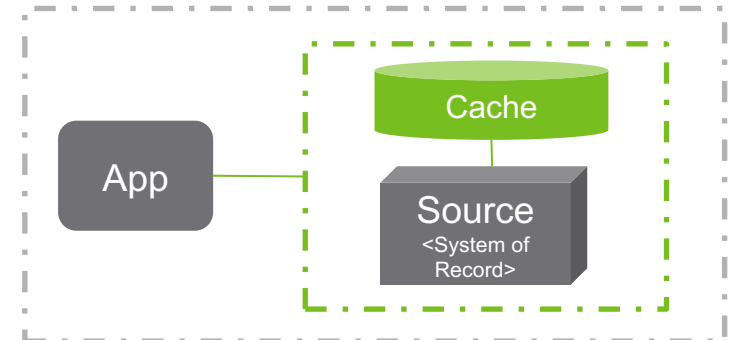


- **NO SQL** data store
 - Fast lookup by key identifiers In-Memory
 - SQL like query (Object Query Language - OQL)
 - Full text-search access
 - Horizontal scalability support
 - High-Availability & Fault Tolerance support
 - WAN replication
 - Triggers/Event notations
 - Stored procedure data processing need
- **Cache** data store
 - API exposed to user interfaces with a real-time interface
 - < 1 second response times
 - Expire cached entries as needed
- Transactional **Operational** data store
 - Persistent with STRONG Consistency – ACID compliance
 - System of record

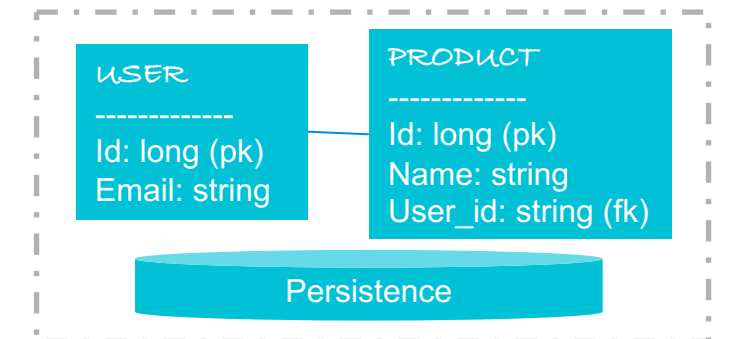
NoSQL Data Store



Cache Data Store



Operational Data Store



GemFire

Fundamentals

Core components

- Data Node – Cache Server – In-memory data storage
- Locator – clients and data nodes controller

Add Data Nodes as needed

- Handle data growth
- Increased processing demands of clients
- Supports resiliency

GemFire cluster

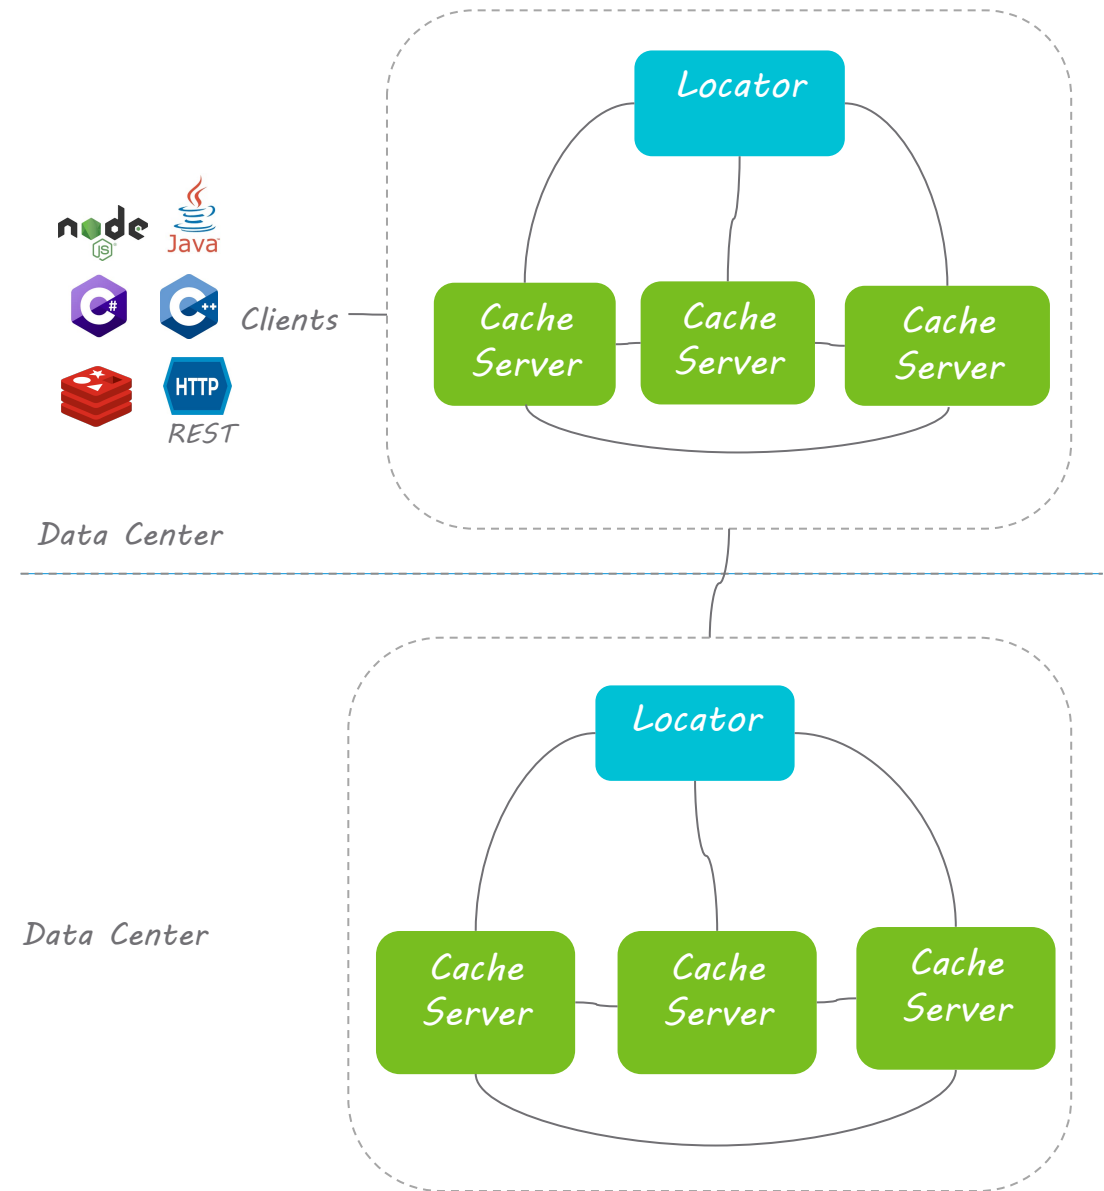
- Connected locators and data nodes

Clients

- Various supported client libraries

WAN Replication

- Replication data across data centers for disaster recovery (DR)
- Active-Active or Active-Passive



Regions

GemFire Region is a database table like data store represented in key/value `java.util.Map` structure

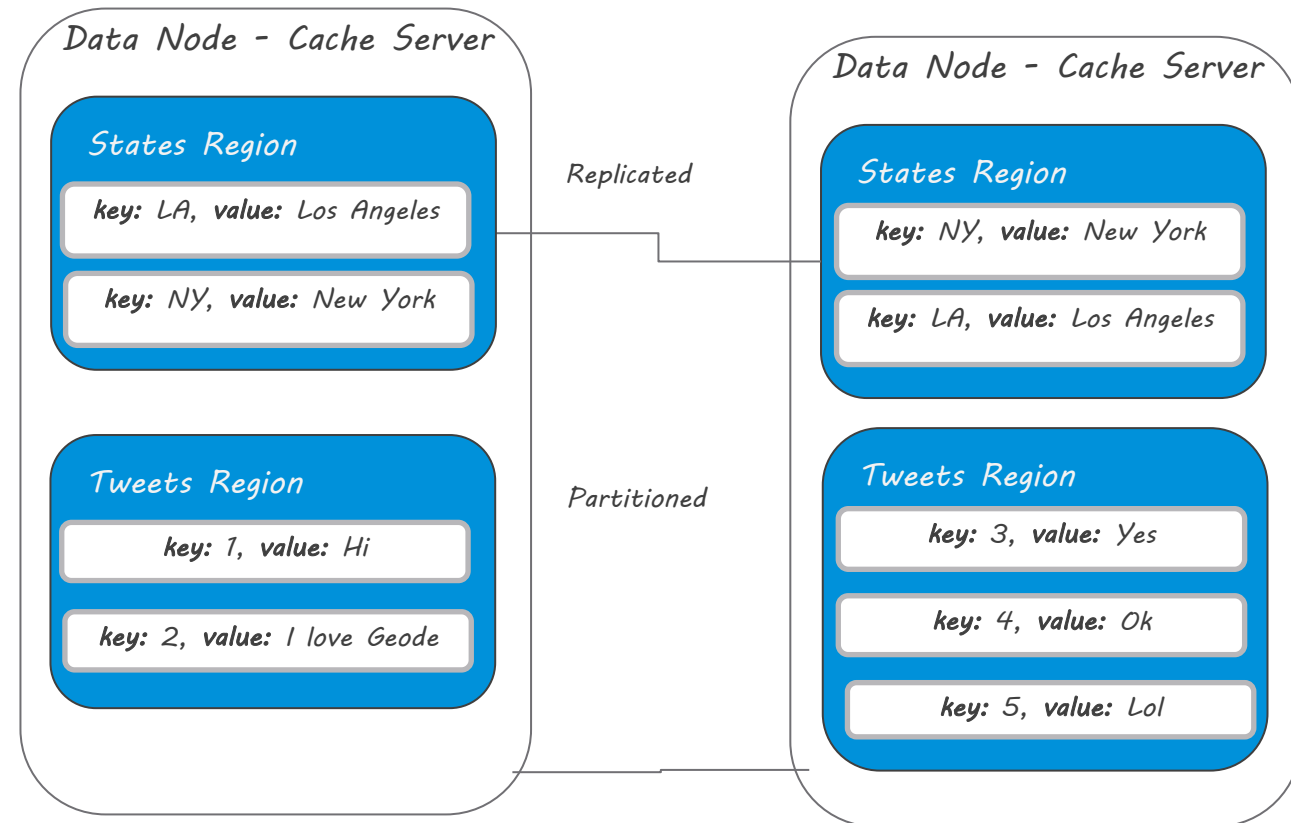
- Region supports querying
 - Ex: select * from /states where code in ('NY', 'LA')
 - Ex: lucene search --regionName=/tweet - queryStrings="*Spring*" --defaultField=tweet
- Events
 - Listeners – triggers data events to client or server-side code
 - Continuous Query – client-side code alerting based on select statements
 - Ex: select * from /tweet where ...
- Transaction support

Cluster Data Policies

- Replicated Region
 - Full copy on each JVM peer
- Partitioned Region
 - Each peer only stores parts of the region contents

```
//Java Code
Region<String,State> region;

region.put(state.code, state);
```



Spring Data Geode

Spring based abstraction layer

- Bootstrapping GemFire
- Spring Data template-based CRUD POJO access, exception translation, transaction management, and query operations.
- Incorporate best practice serialization of managed objects.
- Event driven abstraction using Continuous Query (CQ) to process a stream of events based on interest defined thru the OQL (Object Query Language).

```
@Repository
public interface AccountGeodeRepository extends CrudRepository<Account,String>
{
    Iterable<Account> findByName(String name);

    Iterable<Account> findByNameLike(String name);
}
```

```
@Configuration
@EnableSecurity
@EnableEntityDefinedRegions
@ClientCacheApplication
public class GeodeConfig
{
}
```

```
@Data
@Builder
@NoArgsConstructor
@AllArgsConstructor
@Region
public class Account
{
    private String id;
    private String name;
}
```

```
@Component
class PremiumAccountCqListener {
    private var log = LogManager.getLogger(PremiumAccountCqListener::class.java)

    @ContinuousQuery(name="AccountCq",
        query = "select * from /Account where balance.amount > 100000 "+
            "and (bank_id = 'VMware' or bank_id = 'SPRINGONE')")
    fun handle(cqEvent: CqEvent) {
        var eventOperation = cqEvent.baseOperation
        var key = cqEvent.key

        if(eventOperation.isDestroy) {
            log.warn("Premium Balance Account $key DELETED!!!")
            return
        }

        var newValue = cqEvent.newValue
        log.info("Premium Account $key operation ${eventOperation} executed resulting in $newValue")
    }
}
```


Exercise

GemFire Cluster - Setup

```
package com.vmware.spring.geode.showcase.controller;

import com.vmware.spring.geode.showcase.domain.Account;
import com.vmware.spring.geode.showcase.repositories.AccountRepository;
import lombok.AllArgsConstructor;
import org.springframework.web.bind.annotation.*;

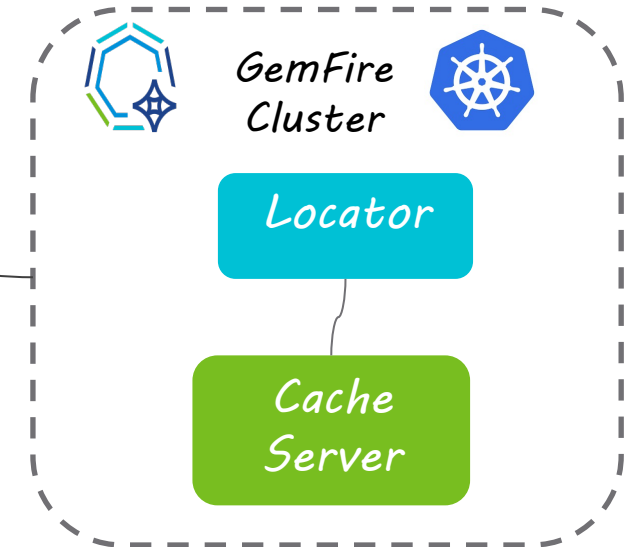
import java.util.Optional;

@AllArgsConstructor
@RestController
public class AccountController
{
    private final AccountRepository accountRepository;

    @PostMapping("save")
    public <S extends Account> S save(@RequestBody S s) { return accountRepository.save(s); }

    @GetMapping("findById")
    public Optional<Account> findById(String s) { return accountRepository.findById(s); }

    @DeleteMapping("deleteById/{id}")
    public void deleteById(@PathVariable String id) { accountRepository.deleteById(id); }
}
```



Account Region
PARTITIONED

```
package com.vmware.spring.geode.showcase.repositories;

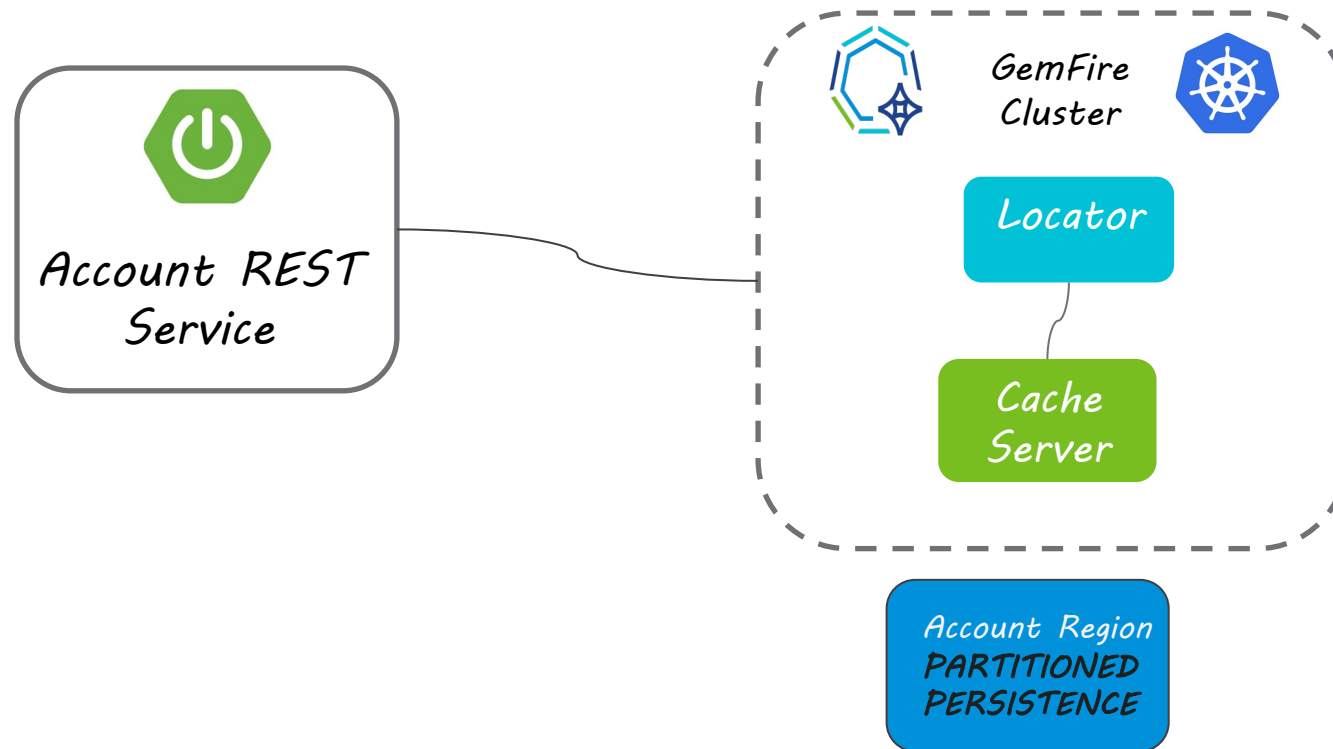
import com.vmware.spring.geode.showcase.domain.Account;
import org.springframework.data.repository.CrudRepository;
import org.springframework.stereotype.Repository;

@Repository
public interface AccountRepository
extends CrudRepository<Account, String>
{
}
```


Exercise

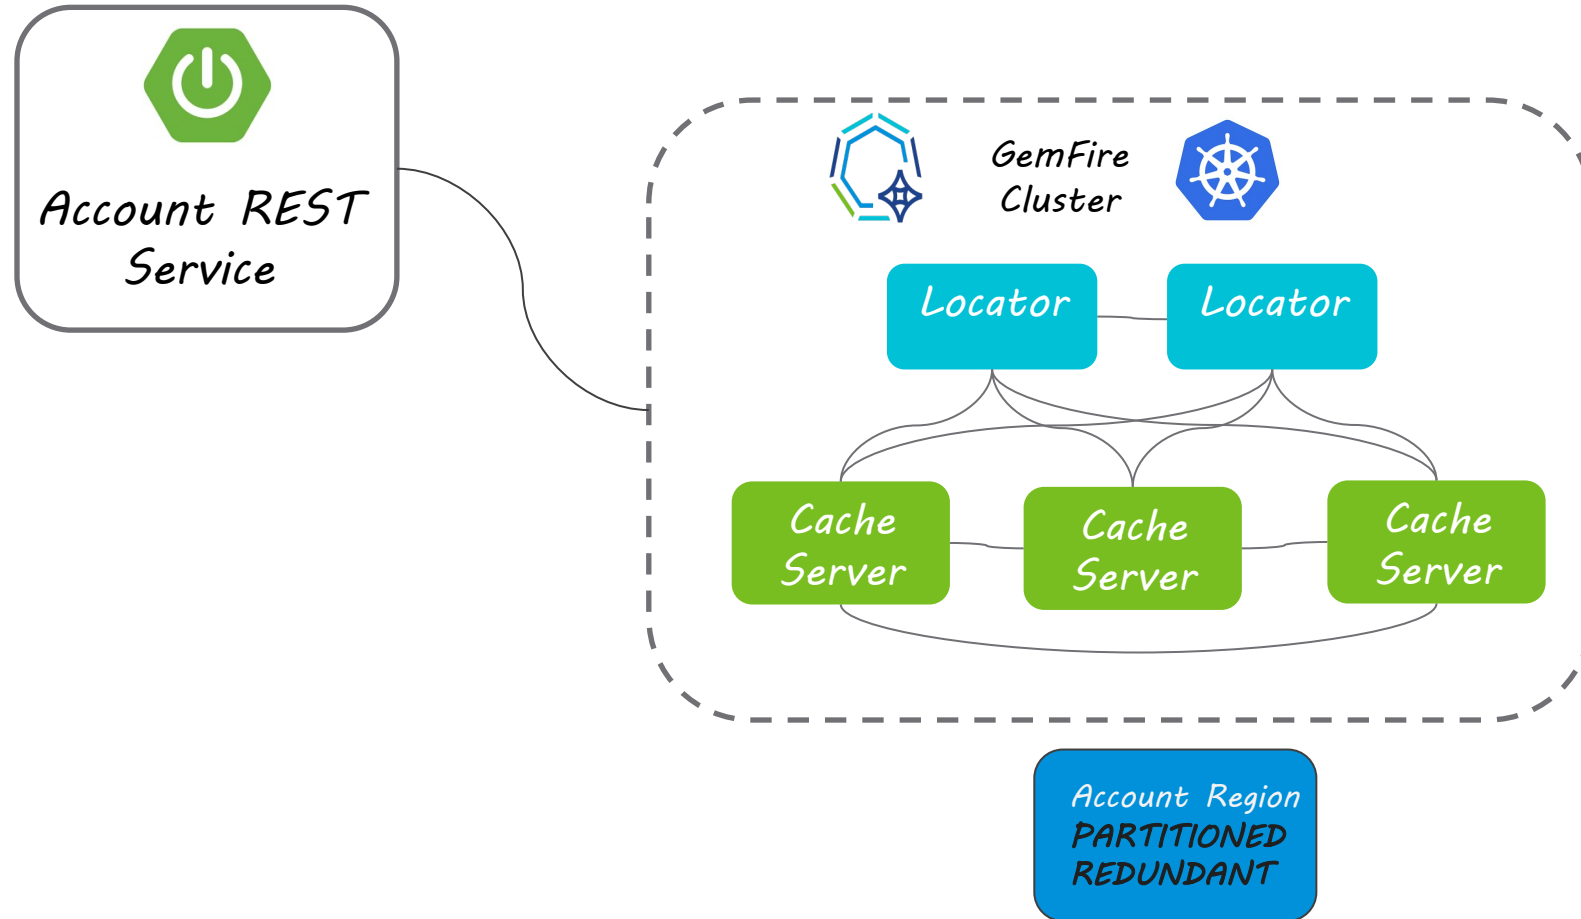
Persistence

<https://github.com/ggreen/spring-geode-showcase.git>



Exercise

Scalability/High Availability



Exercise

Transaction

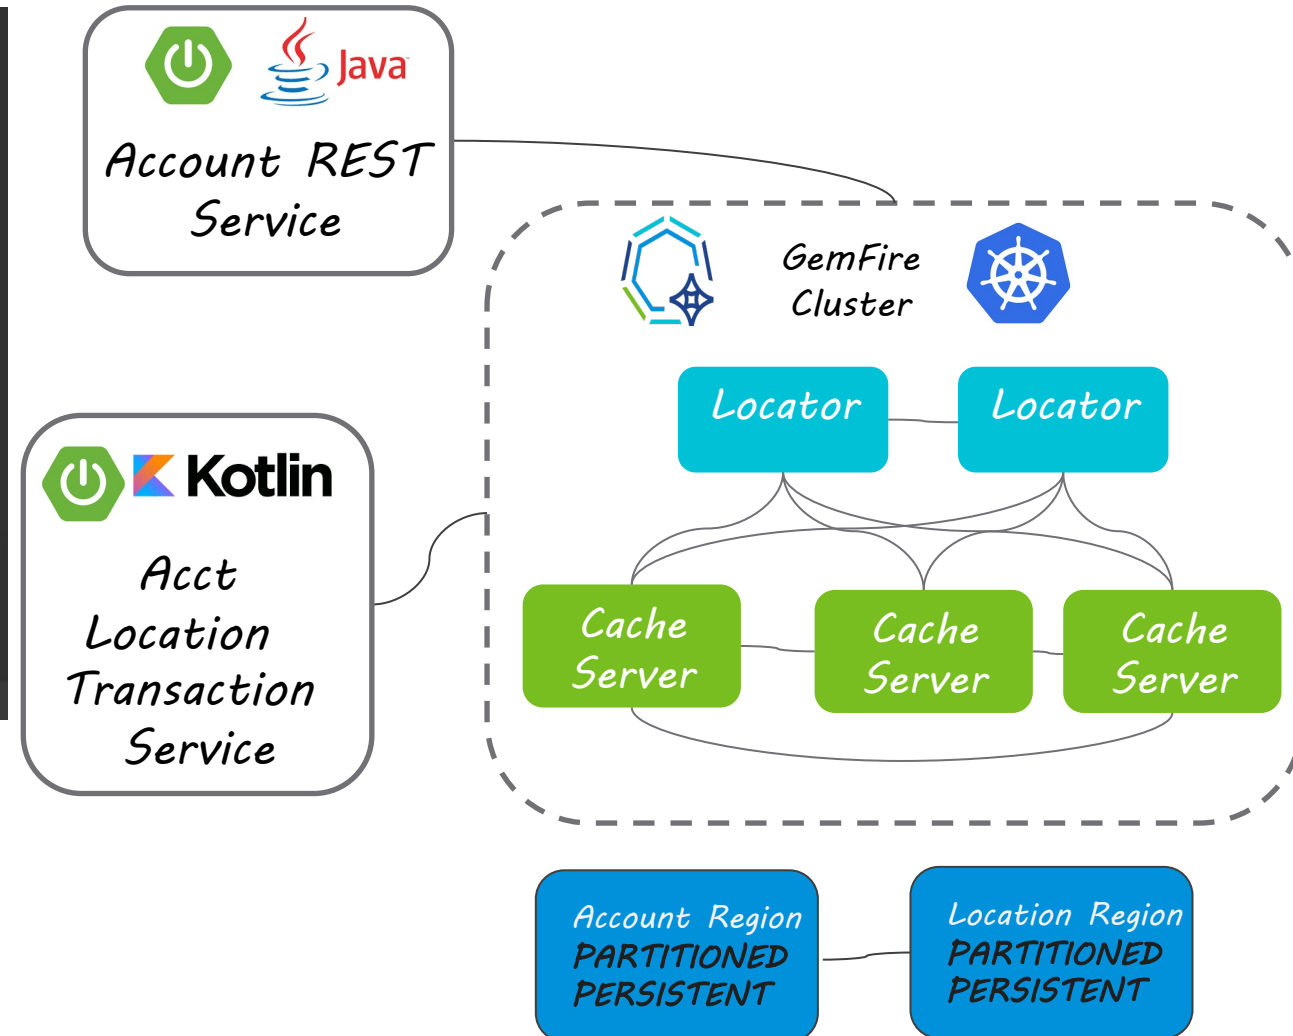
```
@RestController
class AccountLocationController(
    private val accountRepository: AccountRepository,
    private val locationRepository: LocationRepository) {
    private val validZipRegex = "^\\d{5}(?:[-\\s]\\d{4})?\\$".toRegex();

    @PostMapping("save")
    @Transactional
    fun save(@RequestBody accountLocation: AccountLocation) {
        accountRepository.save(accountLocation.account)

        var location = accountLocation.location;
        if(!location.zipCode.matches(validZipRegex))
            throw IllegalArgumentException("Invalid zip code ${location.zipCode}");

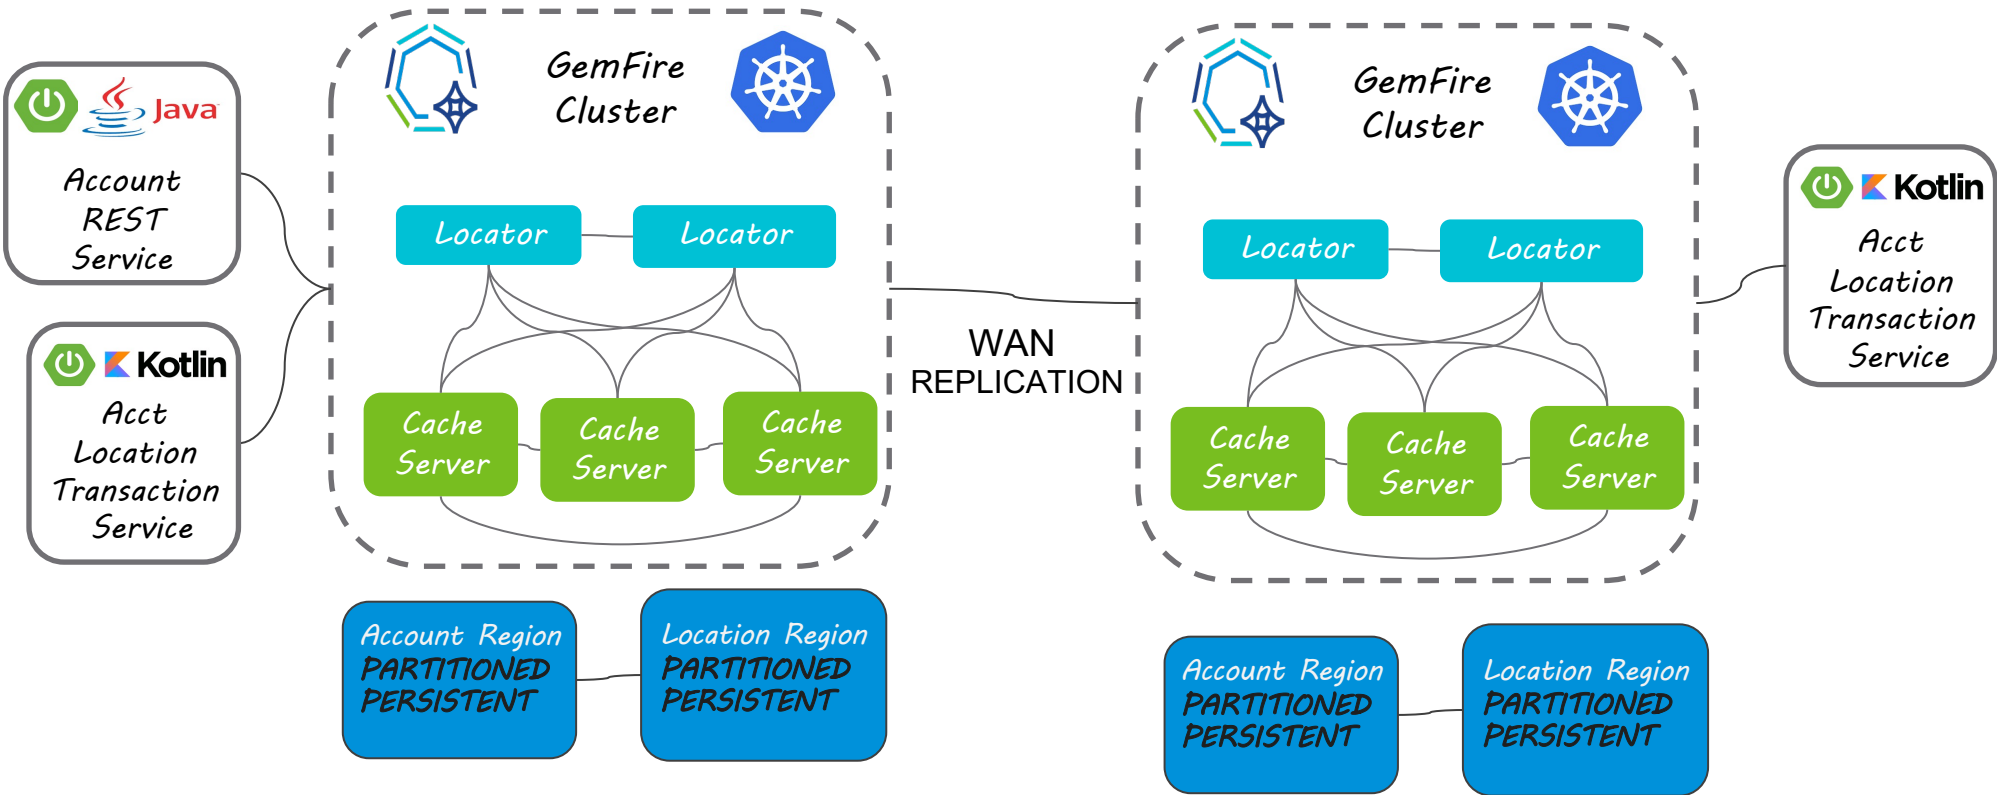
        locationRepository.save(accountLocation.location)
    }
}
```

```
@ClientCacheApplication
@EnableClusterDefinedRegions
@Configuration
@EnableGemfireCacheTransactions
public class GeodeConf
{
}
```



Exercise


WAN Replication



Cacheable

Look aside

- Use Case
 - Data can be loaded as needed
 - Data domains access by id
 - Need to minimize the initial cache storage needs
- Spring Cache Abstraction
 - **@Cacheable** - result is stored into the cache so on subsequent invocations (with the same arguments), the value in the cache is returned without having to execute the method.
 - **@CacheEvict** - perform cache *eviction*, that is methods that act as triggers for removing data from the cache.



```
@Service
class AccountDataService (private val accountRepository : AccountRepository)
: AccountService {
    @CacheEvict(value = ["AccountCache"], key = "#account.id")
    override fun save(account: Account): Account {
        return accountRepository.save(account)
    }

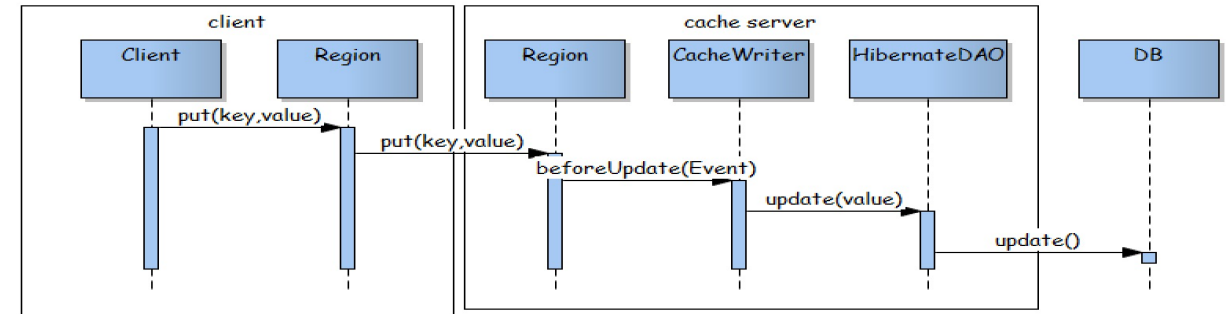
    @Cacheable(value = ["AccountCache"])
    override fun findById(id: String): Account? {
        var optional = accountRepository.findById(id)
        if (optional.isEmpty)
            return null

        return optional.get()
    }
}
```

Cache Writer

Write through

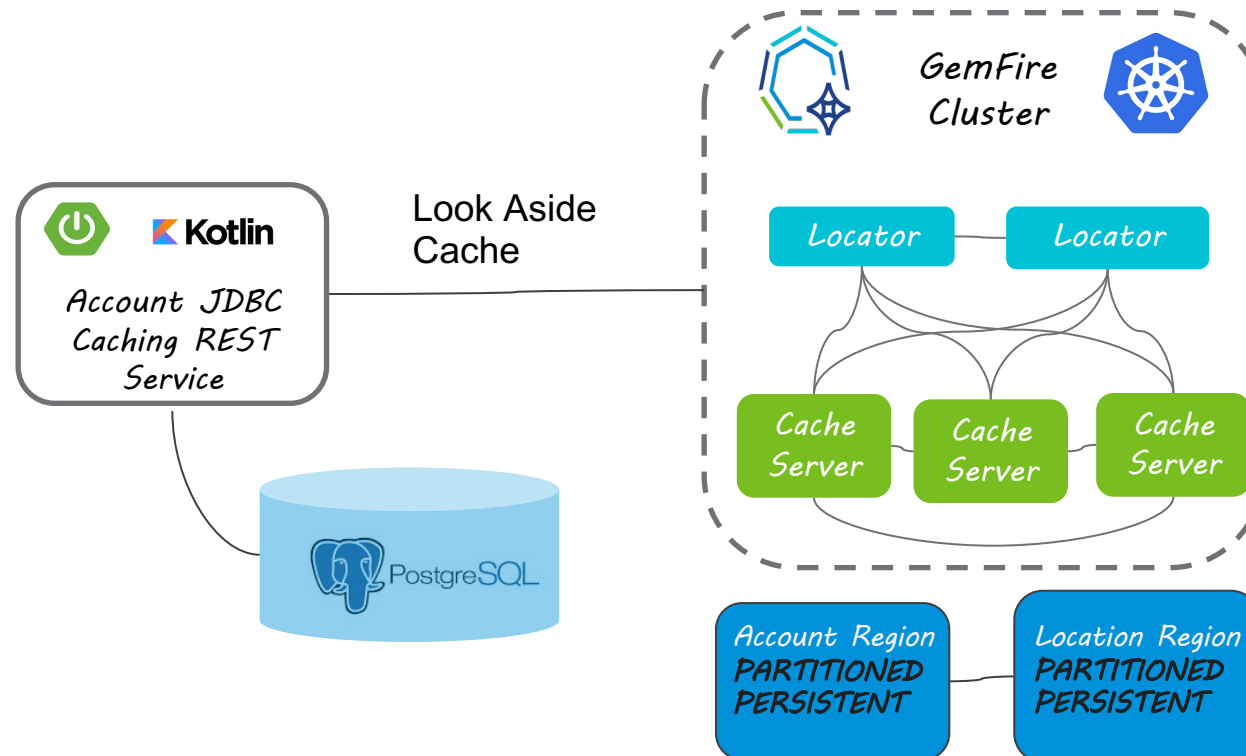
- Use Case
 - Transactional data updates are required to maintain consistency between cache and external data systems
 - Real-time updates are required
 - Update/write rate relatively low
 - Cache write latency is acceptable
- Cache Writer
 - CacheWriters are user-defined and associated with a region.
 - Usage triggered by put key/value entry into a region.
 - CacheWriter's beforeUpdate method may be called in the case of a put of existing record's.
 - Their beforeUpdate, beforeCreate, beforeDestroy, beforeRegionClear, etc. methods are called synchronously before a region or entry in the cache is modified.
 - The region operation may be client and or server-side.



```
class DecisionManagementSystemWriter implements CacheWriter<?, EligibilityDecision> {  
  
    private final DataSource dataSource;  
  
    DecisionManagementSystemWriter(DataSource dataSource) {  
        this.dataSource = dataSource;  
    }  
  
    public void beforeCreate(EntryEvent<?, EligibilityDecision> entryEvent) {  
        // Use configured DataSource to save (e.g. INSERT) the entry to the backend data store  
    }  
  
    public void beforeUpdate(EntryEvent<?, EligibilityDecision> entryEvent) {  
        // Use the configured DataSource to save (e.g. UPDATE or UPSERT) the entry in the  
        backend data store  
    }  
  
    public void beforeDestroy(EntryEvent<?, EligibilityDecision> entryEvent) {  
        // Use the configured DataSource to delete (i.e. DELETE) the entry from the backend data  
        store  
    }  
  
    ...  
}
```


Exercise

Look Aside Cache




Functions

Functions are GemFire's equivalent to database stored procedures

- Execute business logic that is co-located with data in-memory
- Fastest data access patterns
- Functions can be made asynchronous by setting the *hasResult* flag to false and not returning a value.
- Function can be executed programmatically or manually through Gfsh

Execution Types

- **OnRegion** execute on region/partition
 - Executing code in the exact node where a specific key resides in a partitioned Region
- **OnServers** execute on all servers in a pool
 - Executing code simultaneously on all nodes
- **OnServer** execute on a single server in a pool
- **OnMember** execute on a particular server



GemFire Functions


Register with XML

Command line or programmatic API

Execute through command line or programmatic API

```
gfsh> deploy --jar=/tmp/myfunction.jar

gfsh> execute function --id=ClearRegionFunction
--region=/test
```



```
<!-- XML Registration -->
<cache>
  <function>
    <class-
name>demo.UserMgrOnRegionFunction</class-name>
  </function>
</cache>
```

```
Execution execution = FunctionService.onRegion(exampleRegion)
    .withFilter(keysForGet)
    .setArguments(Boolean.TRUE)
    .withCollector(new MyArrayListResultCollector());
```

```
ResultCollector rc = execution.execute(function);
// Retrieve results, if the function returns results
List result = (List)rc.getResult();
```

GemFire Functions Execution

Minimize network hops
with functions

Co-locate data from
different partitioned
regions

- Group related partitioned region data on same member.
 - Improves queries and other operations that access data access time
 - Ex: Co-locate material, types, groups, units of measure, plants, MDM data based of accounts

```
create region --name=Account --type=PARTITION_PERSISTENT
```

```
create region --name=Location --type=PARTITION_PERSISTENT  
--colocated-with=/Account
```

```
public class AccountCountInNyFunction implements Function<PdxInstance>, Declarable  
{  
    private Logger logger = LogManager.getLogger(AccountCountInNyFunction.class);  
    private static final String empty = "";  
    private Cache cache;  
    private QueryService queryService;  
  
    @Override  
    public void execute(FunctionContext<PdxInstance> functionContext)  
    {  
        logger.info(s: "Executing account function");  
  
        ResultSender<String> sender = functionContext.getResultSender();  
  
        if(! (functionContext instanceof RegionFunctionContext)){...}  
        RegionFunctionContext rfc = (RegionFunctionContext) functionContext;  
  
        if(queryService == null)  
            queryService = CacheFactory.getAnyInstance().getQueryService();  
  
        Query query = queryService.newQuery(  
            s: "select count(*) as cnt from /Account a, /Location l where a.id = l.id and l.stateCode = 'NY'");  
    }  
}
```

Exercise

Function execution

```
public class AccountCountInNyFunction implements Function<PdxInstance>, Declarable
{
    private Logger logger = LogManager.getLogger(AccountCountInNyFunction.class);
    private static final String empty = "";
    private Cache cache;
    private QueryService queryService;

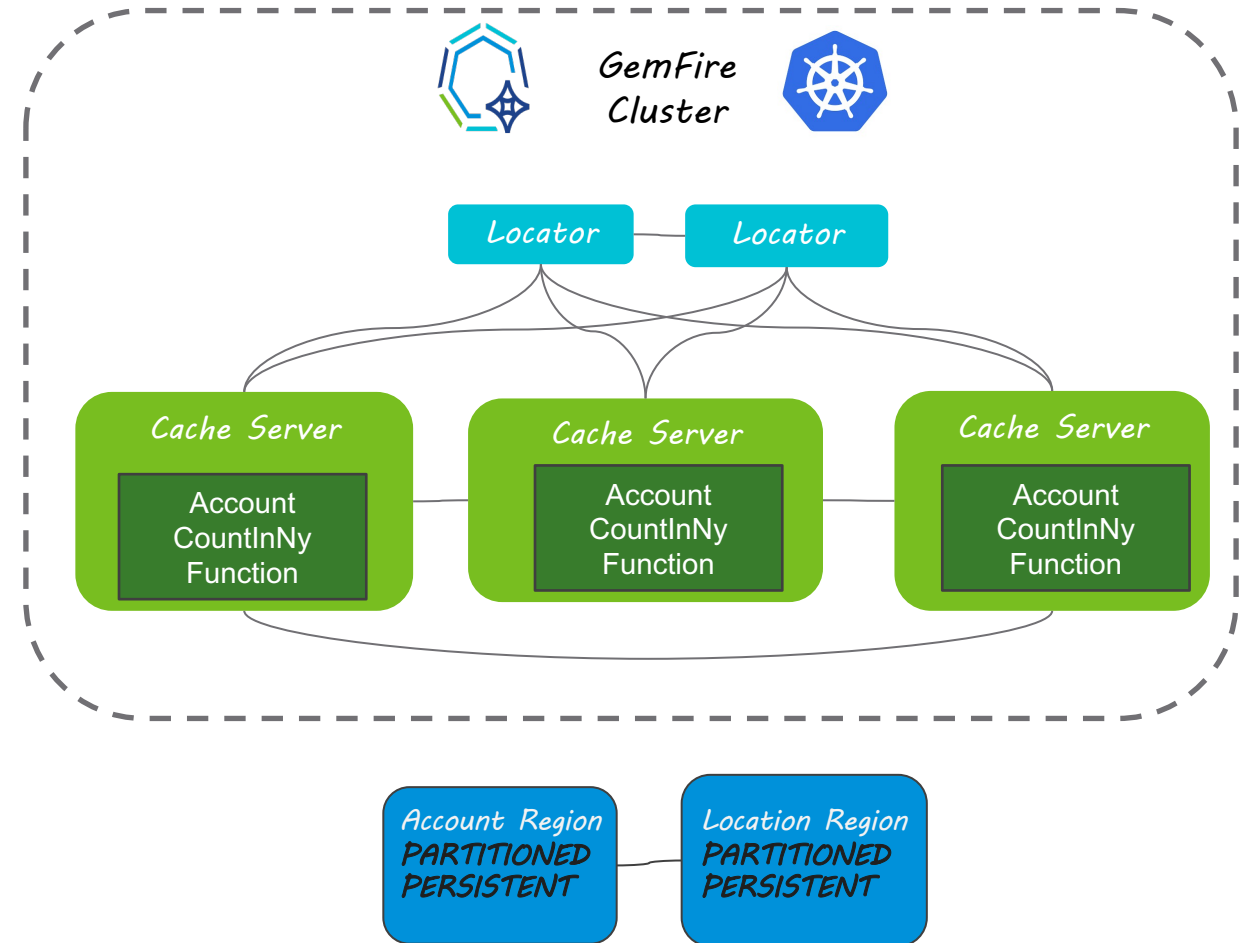
    @Override
    public void execute(FunctionContext<PdxInstance> functionContext)
    {
        logger.info(s: "Executing account function");

        ResultSender<String> sender = functionContext.getResultSender();

        if(! (functionContext instanceof RegionFunctionContext)){...}
        RegionFunctionContext rfc = (RegionFunctionContext) functionContext;

        if(queryService == null)
            queryService = CacheFactory.getAnyInstance().getQueryService();

        Query query = queryService.newQuery(
            s: "select count(*) as cnt from /Account a, /Location l where a.id = l.id and l.stateCode = 'NY'");
    }
}
```



GemFire Development Links

- Core Java/Apache Blog
 - <https://www.baeldung.com/apache-geode>
- Spring Data Geode Blog
 - <https://www.baeldung.com/spring-data-geode>
- Tanzu GemFire Developer Center
 - <https://tanzu.vmware.com/developer/data/tanzu-gemfire/>



Thank You